

# **Oracle Database 11g: OLAP Essentials**

**Student Guide**

D70039GC10

Edition 1.0

December 2008

D57201

**ORACLE®**

**Author**

Brian Pottle

**Technical Contributors  
and Reviewers**

Marty Gubar

Bud Endress

Keith Laker

Kevin Lancaster

Brian Macdonald

Jameson White

**Editors**

Joyce Raftery

Daniel Milne

**Graphic Designer**

Rajiv Chandrabhanu

**Publishers**

Jobi Varghese

Nita Brozowski

**Copyright © 2008, Oracle. All rights reserved.**

**Disclaimer**

This document contains proprietary information and is protected by copyright and other intellectual property laws. You may copy and print this document solely for your own use in an Oracle training course. The document may not be modified or altered in any way. Except where your use constitutes "fair use" under copyright law, you may not use, share, download, upload, copy, print, display, perform, reproduce, publish, license, post, transmit, or distribute this document in whole or in part without the express authorization of Oracle.

The information contained in this document is subject to change without notice. If you find any problems in the document, please report them in writing to: Oracle University, 500 Oracle Parkway, Redwood Shores, California 94065 USA. This document is not warranted to be error-free.

**Restricted Rights Notice**

If this documentation is delivered to the United States Government or anyone using the documentation on behalf of the United States Government, the following notice is applicable:

**U.S. GOVERNMENT RIGHTS**

The U.S. Government's rights to use, modify, reproduce, release, perform, display, or disclose these training materials are restricted by the terms of the applicable Oracle license agreement and/or the applicable U.S. Government contract.

**Trademark Notice**

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

# Contents

## **I Oracle Database 11g: OLAP Essentials**

Agenda 1-2

### **1 Examining the Role of Oracle OLAP Within the Oracle BI/DW Platform**

Objectives 1-2

What Is Business Intelligence? 1-3

The User Dilemma 1-4

The Optimal Information Platform: Security & Compliance 1-5

The Optimal Information Platform: Management & Change 1-6

The Optimal Information Platform: High Availability 1-7

The Optimal Information Platform: VLDB & ILM 1-8

The Optimal Information Platform: DW & Analytics 1-9

The Optimal Information Platform for Business Intelligence 1-10

Physical Layers of a Data Warehouse 1-11

Managing the Data Warehouse 1-12

Oracle OLAP: Part of an Integrated DW Platform 1-13

What Is Oracle OLAP? 1-14

Oracle OLAP Option: 11g Highlights 1-15

Summary Management with Cube-Organized Materialized Views 1-16

DEMONSTRATION: Cube-Organized Materialized Views 1-17

Delivering Rich Analytics to SQL Tools 1-18

Empowering Any SQL-Based Tool 1-19

DEMONSTRATION: Direct SQL Access to Multidimensional Engine 1-20

Summary 1-21

### **2 Examining the OLAP Data Model**

Objectives 2-2

Dimensional Model 2-3

Measures 2-4

Measure Types 2-5

Dimensions 2-6

Example of Dimensions in a Report 2-7

Sharing Dimensions 2-8

Hierarchy 2-9

Hierarchy: Example 2-11

- Level 2-12
- Types of Hierarchy 2-13
- Attributes 2-14
- Attributes and Levels: Examples 2-15
- Dimensional Model Summarized 2-16
- Summary 2-17

### **3 Building an OLAP Cube**

- Objectives 3-2
- What Is an Analytic Workspace? 3-3
- Requirements for Cube Building 3-4
- Steps for Creating a Cube 3-5
- Source Data: OLAPTRAIN Schema 3-6
- Designing the Multidimensional Model 3-7
- What Are Oracle OLAP Cubes? 3-8
- Examining the Multidimensional Model for this Lesson 3-9
- At What Level Should We Load Data? 3-10
- Examining the Multidimensional Model: Dimension Levels 3-11
- Starting AWM 3-12
- Creating a Database Connection 3-13
- Connecting to a Database 3-14
- Creating an Analytic Workspace 3-15
- Analytic Workspace: Creation and Maintenance Privileges 3-16
- Using the AWM Interface 3-17
- Creating OLAP Dimensions 3-18
- Creating a Dimension 3-19
- Dimension Implementation Details 3-20
- Creating Levels 3-21
- Creating Hierarchies 3-22
- Creating Attributes 3-23
- Creating User Attributes 3-24
- Interface for Mapping Dimensions 3-25
- Requirements for Dimension Mapping 3-26
- Mapping Pane Options 3-27
- Table Mapping View: Star Dimension Table 3-28
- Other Dimension Mapping Techniques 3-29
- Using Snowflake Dimension Tables 3-30
- Table Mapping View: Snowflake Dimension Tables 3-31
- Mapping Snowflake Dimensions for Cube MVs 3-32
- Graphical Mapping View 3-33
- Loading Data 3-34

Viewing the Data Load Report	3-35
Using the Data Viewer	3-36
Relational Views of Dimension Data	3-37
Using Templates	3-38
Configuring the Template Source	3-39
Creating Dimensions Using Templates	3-40
Practice 3-1: Overview	3-41
Creating OLAP Cubes	3-42
Creating Cubes	3-43
Cubes: Select the Dimensions	3-44
Cubes: Select the Storage Options	3-45
Cubes: Select the Partitioning Options	3-46
Cubes: Select the Aggregation Rules	3-47
Cubes: Select the Precompute Options	3-48
Creating Measures	3-49
Interface for Mapping Cubes	3-50
Mapping Cubes to Data Sources	3-51
Loading and Aggregating Data	3-52
Loading Data: Specify Target Objects	3-53
Loading Data: Specify a Job Schedule	3-54
Viewing the Cube Build Log	3-55
Monitoring the Build Using CUBE_BUILD_LOG	3-56
Viewing Measures	3-57
Relational Views of Measure Data	3-58
Reviewing Database-Generated Objects	3-59
Summary	3-60
Practice 3-2: Overview	3-61

## **4 Examining Cube-Organized Materialized Views**

Objectives	4-2
Materialized Views: Rationale	4-3
Materialized Views: Typical MV Architecture	4-4
Materialized Views: Challenges in Ad Hoc Query Environments	4-5
Cube-Organized Materialized Views	4-6
Implementing Cube MVs	4-7
Implementing Cube MVs: Design	4-8
Implementing Cube MVs: Enable	4-9
How Cube MVs Work	4-11
MV Refresh Options	4-12
MV Refresh Modes	4-13
MV Refresh Methods	4-14

- Fast Solve 4-16
- Implementing Cube MVs: Build Cube 4-17
- Cube MV Objects Created by the Build 4-18
- Implementing Cube MVs: Prepare Schema 4-19
- Relational Schema Advisor 4-21
- Obtaining Query Rewrite Results 4-22
- Summary 4-23
- Practice 4: Overview 4-24

## **5 Creating Calculated Measures**

- Objectives 5-2
- Agenda 5-3
- Types of OLAP Calculations 5-4
- Calculated Measures 5-5
- Calculated Measures and Inherited Dimensionality 5-7
- Characteristics of Calculated Measures 5-8
- Calculated Measures Example 5-10
- Aggregations 5-11
- Aggregations Can Vary by Dimension 5-12
- Agenda 5-13
- Creating Common Business Calculations 5-14
- Using the Calculation Builder 5-15
- Time Series Calculations 5-16
- Defining Time Series Calculations: Parallel Period 5-17
- Explanation of Parallel Period 5-18
- Parallel Period Example 5-20
- Defining Time Series Calculations: Period-to-Date 5-21
- Explanation of Period-to-Date 5-22
- Period-To-Date Example 5-23
- Defining a Hierarchy Calculation: Share 5-24
- Share Calculations Leverage Hierarchy “Parents” 5-25
- Share Example 5-26
- Nesting Analytic Calculations 5-27
- Agenda 5-28
- Overview of OLAP Calculation Expression Syntax 5-29
- Similarity to Standard SQL Syntax 5-30
- Syntax Is Modeled on SQL Functions 5-31
- Example: Using the RANK Function 5-32
- Example: Using the Moving Average Function 5-33
- Modifying Existing Expressions 5-34
- Creating Custom Expressions 5-35

Viewing Custom Calculation Results	5-36
Referencing Dimension Members in Custom Expressions	5-37
Using OLAP DML Code in Calculations	5-38
Introduction to OLAP DML	5-39
Referencing Objects in the OLAP DML	5-40
Using the OLAP Worksheet	5-41
Using OLAP DML in Calculated Measures	5-42
Using OLAP DML: An Example	5-43
Examining the OLAP DML Program	5-44
Calling a DML Program in a Calculation Expression	5-45
Agenda	5-46
Creating Calculations from an XML Template	5-47
Resulting Calculations from XML Templates	5-48
Summary of OLAP Calculation Benefits	5-49
Summary	5-50
Practice 5: Overview	5-51

## **6 Using SQL to Query Oracle OLAP Cubes**

Objectives	6-2
Exposing OLAP Cubes Using Relational Views	6-3
Comparing “Star” Models	6-4
Identifying Cube and Dimension Views	6-5
Cube and Dimension Views in the Oracle Data Dictionary	6-6
Managing Cube Views	6-7
Dimension Views	6-8
Dimension Hierarchy View Example	6-9
Cube Views	6-10
One Cube View for All Summaries	6-11
Leveraging Cube Summaries	6-12
Simplifying Analytic Queries	6-13
Cube Queries Are Similar to Star Queries	6-14
Join Optimization	6-15
Querying Cubes: The Basics	6-16
Easily Query Multiple Summary Levels	6-17
Hierarchy Definition Is Part of the Model: Drilling	6-18
Hierarchy Definition Is Part of the Model: “Family” Queries	6-19
Level-Based Queries	6-20
Using Level Filters	6-21
Joining Cube Views to Tables	6-22
Aggregating Over Attributes	6-23
SQL Processing for OLAP Data	6-24

Aggregating Over Attributes 6-25  
Using the CUBE\_TABLE Function 6-27  
Using CUBE\_TABLE for Cube Queries 6-28  
Using CUBE\_TABLE for Dimension Queries 6-29  
Using CUBE\_TABLE: Cube Query for Specific Hierarchies and Measures 6-30  
Summary 6-31  
Practice 6: Overview 6-32

## **7 Enhancing Analytic Content**

Objectives 7-2  
Forecasting in OLAP 7-3  
Forecasting Scenario 7-4  
Creating the Forecast 7-5  
Create a Forecast Cube 7-6  
Add the Forecast Measures 7-7  
Examine the Forecast Program 7-8  
Create Custom Cube Script 7-9  
Run the Cube Script to Create Forecast Data 7-10  
Understanding Crossover Calculations 7-11  
Viewing Forecast Data 7-13  
Summary 7-14  
Practice 7: Overview 7-15

## **8 Using Ad Hoc Query and Reporting Tools Against OLAP Data**

Objectives 8-2  
Agenda 8-3  
What Is Oracle Application Express? 8-4  
What Is an APEX Workspace? 8-5  
Defining Roles 8-6  
Logging In to the Workspace 8-7  
APEX Components for Developers 8-8  
Using the SQL Commands Tool 8-9  
Running SQL Commands 8-10  
Accessing Application Builder 8-11  
Creating an Application 8-12  
Creating a Simple Interactive Report 8-13  
Running the Application 8-14  
Using an Interactive Report 8-15  
Other Interactive Reporting Options 8-16  
Examining an Exploratory Report 8-17



Using the Developer Toolbar	8-18
Viewing Application Page Components	8-19
Enabling Ad Hoc Exploration	8-20
Methodology for Enabling Ad Hoc Navigation	8-21
Defining Parameters for Drilldown	8-22
Single Parameterized Query Drives Analysis	8-23
Drilldown Results	8-24
Using APEX Report Links	8-25
Demonstration 1: Using APEX with OLAP Data	8-26
What Is Oracle BI Answers?	8-27
BI Tool Metadata Administration	8-28
Standard BI EE Metadata Overview	8-29
Mapping Physical to Business Model	8-30
Additional BI EE Metadata Administration for OLAP	8-31
Add Logical Table Source Per Level: Apply “Level” Condition	8-32
Add Logical Table Source Per Level: Table Is Source for Multiple Levels	8-33
Automatically Add Level Conditions	8-34
Demonstration 2: Using BI Answers with OLAP Data	8-35
Summary	8-36
Practice 8: Overview	8-37

## **9 Implementing Cube Security**

Objectives	9-2
Security Terms	9-3
Authentication	9-4
User Privileges	9-5
User Privileges for MV Creation	9-6
Authorization Requirements	9-7
Grant Access to Cube Objects	9-9
Source Tables and Views	9-10
Requirement: More Granular Control	9-11
Security Independent of Aggregation Rules	9-12
Alternative Aggregation Method	9-13
Aggregation, Then Filter	9-14
Security Scenario	9-15
Securing Access to Data	9-16
Data Security Policy in AWM	9-17
Creating a Security Policy	9-18
Specifying Security Rules	9-19
Viewing Query Results	9-20
Securing Access to Data	9-21

Virtual Private Database 9-22  
Defining a VPD Security Policy 9-23  
Creating a Policy Function 9-24  
Security Method Benefits 9-25  
Which Security Method Should I Use? 9-26  
Summary 9-27  
Practice 9-1 Overview 9-28

## **10 Designing Cubes for Performance and Scalability**

Objectives 10-2  
Goals of Cube Design 10-3  
Designing for Performance and Scalability 10-4  
Comparing Tables and Cubes 10-6  
Common Features of Tables and Cubes 10-7  
Shared Purpose of Common Features 10-8  
Implementing Common Features 10-9  
Cube Design Roadmap 10-11  
Remaining Lesson Agenda 10-12  
Agenda: Constraints 10-13  
Understanding Cube Storage 10-14  
Comparing Fact Tables and Cubes 10-15  
Examining Cube Storage 10-16  
Cube Storage Example 10-17  
Agenda: Indexes 10-18  
Examining Dimensions 10-19  
Understanding Data Blocks 10-20  
Examining Data Pages 10-21  
Pages and Array-Based Data Storage 10-22  
Examining Array-Based Data Storage 10-23  
Array-Based Storage and Data Sparsity 10-24  
Introducing the Composite Index 10-25  
Candidate for a Composite Index 10-26  
Structure of a Composite Index 10-27  
Cube Storage with a Composite Index 10-28  
Impact of Composite Indexes on Pages 10-29  
Impact of Composite Indexes on Performance 10-30  
Summarizing Composite Index Concepts 10-31  
Creating Composite Indexes 10-32  
Types of Composite Indexes 10-33  
Uncompressed Composite Index 10-34  
Uncompressed Composite Index Example 10-35

Examining Sparsity in the Example	10-36
Compressed Composite Index	10-37
Compressed Composite Index Example	10-38
Impact of the Compressed Composite Index	10-39
Benefits of Compression	10-40
Summarizing Data Sparsity Characteristics	10-41
Creating a Compressed Composite Index	10-42
Composite Indexes and Dimension Order	10-43
Composite Index Recommendations	10-44
Agenda: Partitioning	10-45
Cube and Table Partitioning Are Similar	10-46
Understanding Partitioned Cubes	10-47
Partitioned Cube Types	10-48
Balancing Partitioned Cubes	10-49
Partitioned Cubes: Example A	10-50
Partitioned Cubes: Example B	10-51
Local Versus Global Composites	10-52
Local and Global Composite Considerations	10-53
Partitioned Cubes and Incremental Update	10-54
Partition Pruning	10-55
Practical Recommendations for Partitioning on Time	10-56
Recommendation Strategy for Partitioned Cubes	10-57
Agenda: Aggregation	10-58
Precompute Settings	10-59
Aggregation Settings and Performance Considerations	10-60
Aggregation Operators	10-61
Certain Operators Are Compressible	10-62
Certain Operators Do Not Compress Well	10-63
Summary of Aggregation Operators and Compression	10-64
Aggregation Hierarchies	10-65
Aggregation Hierarchies: Recommendations	10-66
Aggregation Methods	10-67
Understanding Cost-Based Aggregation	10-68
Using Cost-Based Aggregation	10-69
Understanding Level-Based Aggregation	10-70
Specifying Level-Based Aggregation	10-71
Aggregation Methods: Recommendations	10-72
Agenda: Dimensioning the Cube	10-73
Dimensioning the Cube	10-74
Dimensioning the Cube: An Example	10-75
Dimensioning Options	10-76

Dimensioning Cube for MV Rewrite 10-77  
Summary of Recommendations 10-78  
Summary 10-79

## **11 Performance Tuning**

Objectives 11-2  
Cube Storage 11-3  
Cube Building Tasks 11-4  
Loading Data from Sources: Part 1 11-5  
Loading Data from Sources: Part 2 11-6  
Loading Data: Things to Consider 11-7  
Cube Building Tasks 11-8  
Aggregating Data 11-9  
Aggregating Data: Things to Consider 11-10  
Cube Building Tasks 11-11  
Updating the Cube 11-12  
Updating the Data: Considerations 11-13  
Road to Success 11-14  
Balanced Configurations: Solving Hardware Resource Issues 11-15  
Unbalanced Systems Lead to Performance Bottlenecks 11-16  
What Is a Balanced Unit? 11-17  
Sample Balanced Unit 11-18  
Full Range of DW Solution Options 11-19  
Cube Partitioning 11-21  
Parallelism 11-22  
Partitioning and Aggregation 11-23  
Example: Impact of Partitioning on Aggregation Performance 11-24  
Managing Performance Issues 11-25  
Automatic Workload Repository 11-26  
Automatic Database Diagnostic Monitor (ADDM) 11-27  
Example: Why Is the Data Load Slow? 11-29  
Using ADDM to Diagnose the Problem 11-30  
What Were the Performance Hotspots? 11-31  
Review ADDM Recommendations 11-32  
Findings Detail 11-33  
Re-Execute Build and Examine Results 11-34  
Summary 11-35

## **Appendix A: Practices and Solutions**

## **Appendix B: Troubleshooting Query Rewrite for Cube-Organized Materialized Views**